

Privacy-Preserving Machine Learning

Exploring Homomorphic Encryption for Medical Diagnosis

Matthew Townsend

CS540 AI Project

Dataset

- **Source:** UCI Diabetes Dataset (`diabetes.csv`)
- **Task:** Binary classification (Diabetes presence vs. absence)
- **Features:** 8 medical attributes (glucose, blood pressure, BMI, etc.)
- **Samples:** ~768 total (614 train, 154 test at 80/20 split)
- **Preprocessing:** StandardScaler normalization

Plaintext Model Training

```
lr = LogisticRegression(max_iter=1000, random_state=13)
svm = SVC(kernel='linear', random_state=13)
```

- Trains two linear models on plaintext data
- **Logistic Regression:** Probabilistic classifier with learned weights
- **SVM (Linear Kernel):** Maximum margin classifier
- Both are compatible with HE because they only require dot products + bias (no non-linear activations)

Why Linear Models? Non-linear activation functions (ReLU, sigmoid) are expensive to compute in the encrypted domain and introduce significant noise/performance penalties.

Homomorphic Encryption (HE)

Allows computation on encrypted data **without decrypting it**:

- Plaintext: $z = a \cdot x + b$
- Encrypted: $\text{Enc}(z) = \text{Enc}(a) \cdot \text{Enc}(x) + \text{Enc}(b)$ ✓ Same result when decrypted

CKKS Scheme

- **Approximate arithmetic** (slight rounding errors)
- **Supports addition and multiplication** on encrypted numbers
- **Scales by 2^p** to handle fractional arithmetic
- **Noise grows** with each operation; eventually swallows signal

CKKS (8192) - High Precision

```
poly_mod_degree_8192 = 8192  
coeff_mod_bit_sizes_8192 = [60, 40, 40, 60]  
enc_training_ckks_8192.global_scale = 2 ** 40
```

- **Polynomial Degree:** Larger N = higher capacity and precision, but slower computation
- **Coefficient Moduli:** Larger bit sizes provide more "room" for numbers before overflow
- **Global Scale:** Fractional precision in encrypted arithmetic (~1 trillion scale factor)

CKKS (4096) - Lower Precision

```
poly_mod_degree_4096 = 4096  
coeff_mod_bit_sizes_4096 = [40, 21, 21, 21, 21, 40]  
enc_training_ckks_4096.global_scale = 2 ** 20
```

- **Polynomial Degree:** Half the capacity of 8192, faster but noisier
- **Smaller coefficient moduli:** Less precision available
- **Lower scale:** Reduced fractional precision (~ 1 million scale factor)

Purpose: This creates an intentional "precision gap" to demonstrate how HE performance degrades with reduced security margins.

Threshold Calibration

```
threshold_lr_ckks = np.mean(y_lin_enc_train_lr)
```

- Computes the encrypted predictions on 50 training samples
- Takes the mean as the decision boundary
- In HE, decryption can introduce minor noise. We calibrate the threshold to the encrypted domain to ensure fair comparison.

Homomorphic Inference

For each model (LR/SVM) and context (8192/4096):

```
sample_enc = ts.ckks_vector(enc_training_ckks_8192, X_test_scaled[i])
y_lin_enc = sample_enc.dot(lr_weights) + lr_bias    # Encrypted dot product
y_lin_dec = y_lin_enc.decrypt()                    # Decrypt result
y_pred = 1 if y_lin_dec[0] > threshold else 0      # Thresholded classification
```


- **ts.ckks_vector()**: Encrypts plaintext feature vector
- **.dot()**: Computes encrypted dot product (weights · features)
- **decrypt()**: Returns ciphertext to plaintext (only private key can do this)
- All operations are mathematically equivalent to plaintext computation, but with added noise



Learning Curves

Iterates over 5 training set sizes (20%, 40%, 60%, 80%, 100%):

1. **Trains new models** on each subset
2. **Calibrates new thresholds** for both CKKS contexts (separate calibration per subset)
3. **Evaluates all 6 combinations** on the full test set
4. **Stores F1 scores** to track how performance changes with training data volume

Key Insight: The learning curves reveal whether HE models can improve with more training data or if they're fundamentally limited by encryption noise.





Performance Gap Analysis

Plaintext vs. CKKS performance gaps indicate:

- **8192 > 4096**: Higher polynomial degree reduces noise
- **Small gap (<5%)**: HE viable for this task
- **Large gap (>20%)**: Noise floor dominates; dataset/model not suitable for HE

Timing Comparison

- Plaintext: <1 millisecond per sample
- HE: 500ms - 1s per sample
- Speedup needed: ~1000-10000x to make HE practical for real-time inference

Conclusion

This project demonstrates the **fundamental tradeoff in HE**:

- **Security**: Data stays encrypted; server never sees plaintext
- **Accuracy**: Noise from encryption reduces model performance
- **Speed**: Encrypted operations 1000x+ slower than plaintext

HE is ideal for **privacy-critical, latency-tolerant** applications (e.g., outsourced analysis of sensitive medical data). For real-time inference, traditional encryption (TLS) + secure enclaves remain superior.

References

- **TenSEAL:** <https://github.com/OpenMined/TenSEAL>
- **CKKS Scheme:** Cheon et al. (2017) "Homomorphic Encryption for Arithmetic of Approximate Numbers"
- **Diabetes Dataset:** UCI ML Repository